Raters

For raters, consider the average note factor of notes they rate helpful and not helpful, along with how many notes with positive/negative factors they rated helpful/not helpful?

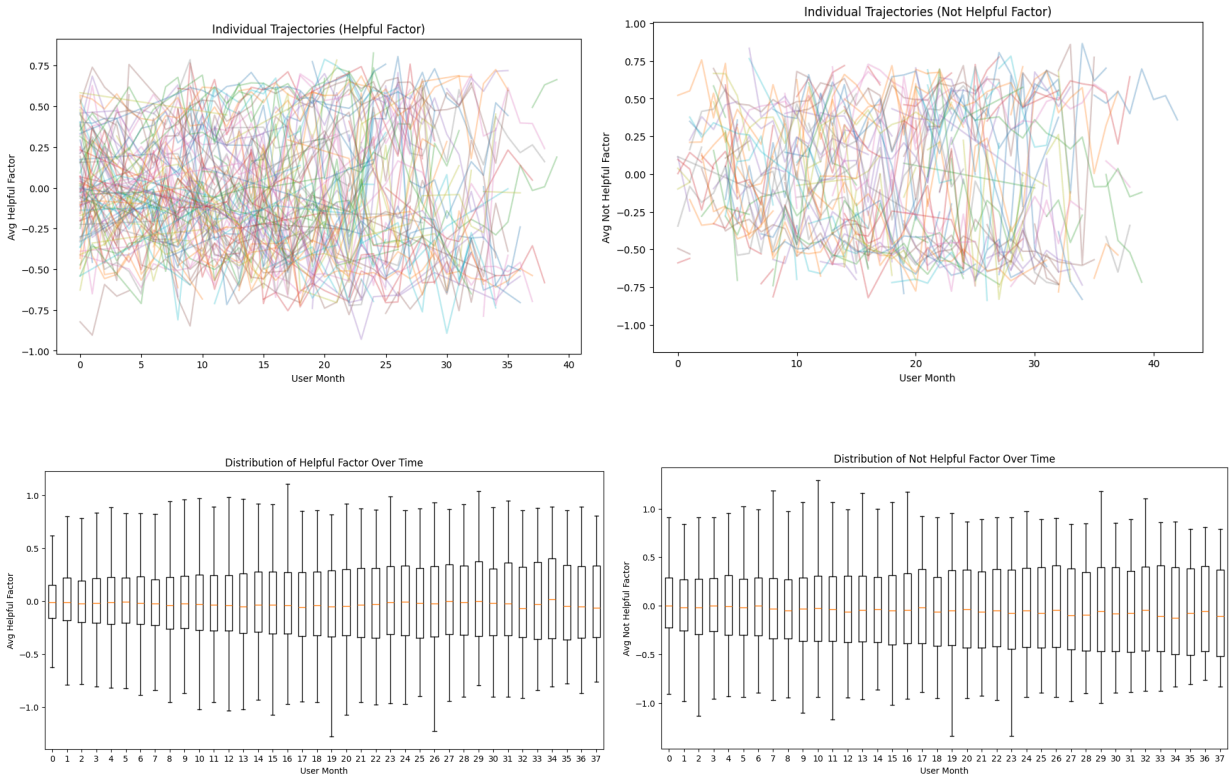

Individual trajectories of avgHelpfulFactor and avgNotHelpfulFactor appear highly noisy and do not exhibit a consistent upward or downward trend. This is supported by the population-level box plots, which show that the distribution of avgHelpfulFactor and avgNotHelpfulFactor remains centered around zero with a stable spread across user months. Together, these results suggest that raters do not get better or worse at picking sides over time, and they stay pretty different from each other throughout

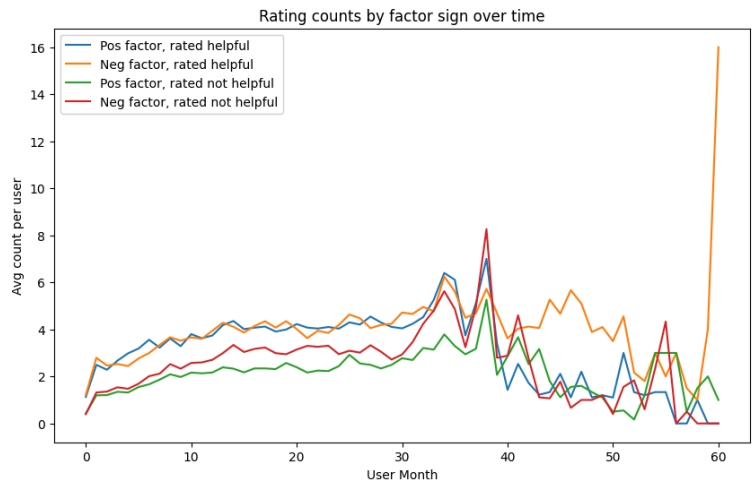

Helpful ratings (blue and orange) are consistently higher than not-helpful ratings (green and red), meaning raters are more likely to call notes helpful regardless of factor sign. Negative-factor notes rated helpful (orange) are slightly above positive-factor notes rated helpful (blue) for most of the early period, suggesting a mild tendency to rate left-leaning notes as helpful, but the difference is very small. The two not-helpful lines track each other closely and stay low throughout, meaning raters don't strongly distinguish between positive and negative factor notes when rejecting them. It’s interesting how they follow a similar trend up until month 40, especially the spike near month 38.

For raters, also consider the percent of ratings they gave that were correct, along with the average note intercept of notes they rated helpful and not helpful.

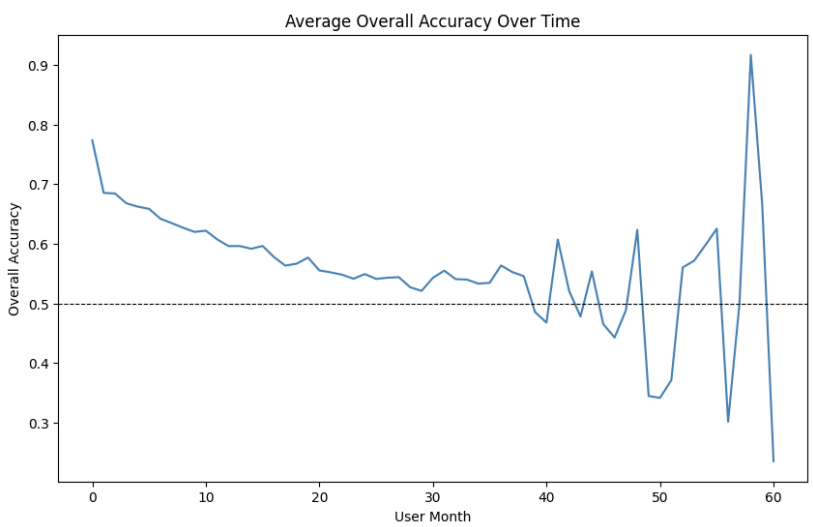

Overall accuracy decreases from about 0.78 at month 0 to roughly 0.55 by month 35, after which estimates become unreliable due to sparse samples. This was unexpected since you would expect raters to get better over time, but the opposite happened. One

possible explanation is that early on, new users are probably shown notes that already have a lot of ratings and a clear outcome, so they are easier to get right. Over time they are shown newer, less-rated notes where the answer is less obvious. The other possibility is that long-term users tend to go looking for more controversial notes, which are harder to judge correctly.

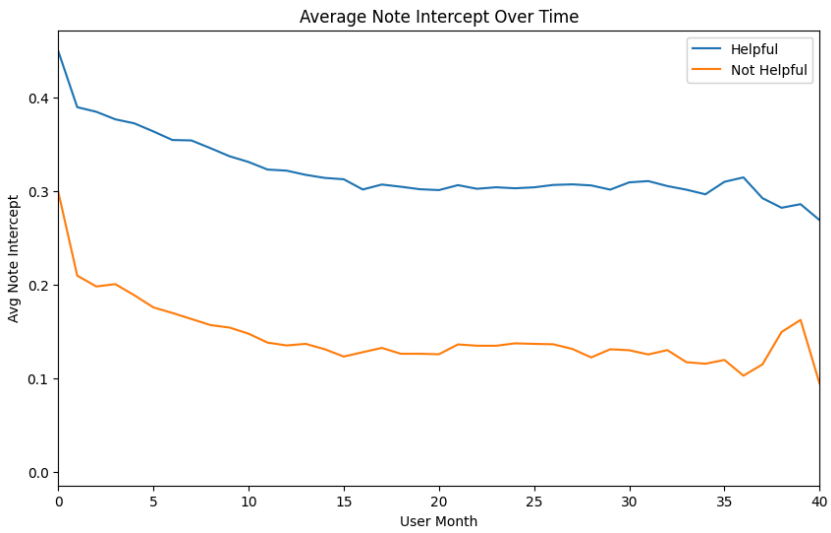

Average note intercept for notes that were rated helpful, were consistently higher compared to average note intercept for notes that were rated not helpful. It is interesting that the overall trend, along with the gap between the two lines stay stable throughout the entire time, where both lines drop in the first few months, which is similar to what we saw in overall accuracy, where early notes tend to have more clearer outcomes. After around 15 months, they start to flatten out.

So overall, raters maintain a consistent ability to distinguish between higher and lower quality notes over time, but their raw accuracy declines. This suggests that raters are not getting worse at judging notes, but rather that the notes they encounter become harder to judge correctly as time goes on.

Writers

For writers, consider the average note factor of notes the user wrote

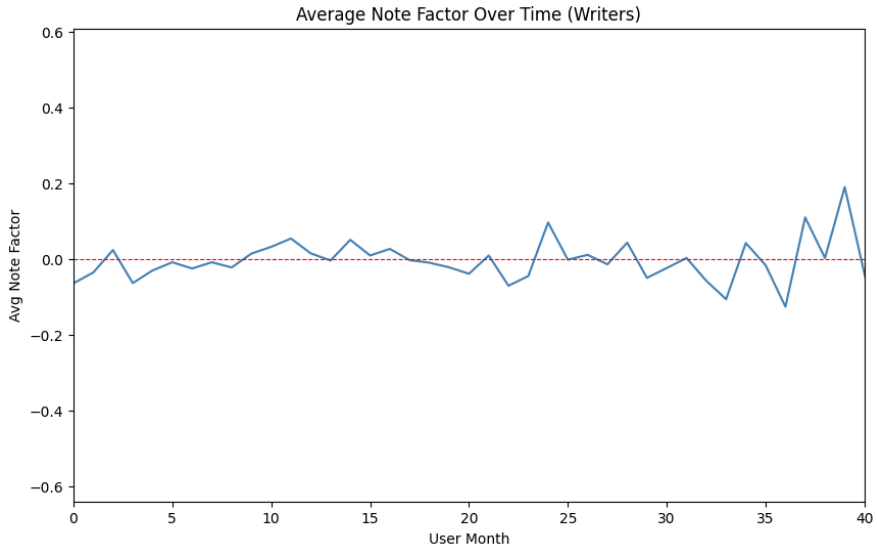

The average note factor over time stays consistent throughout, hovering around zero with no clear upward or downward trend. This suggests that on average, writers do not have a tendency to become politically biased, and remain relatively neutral.

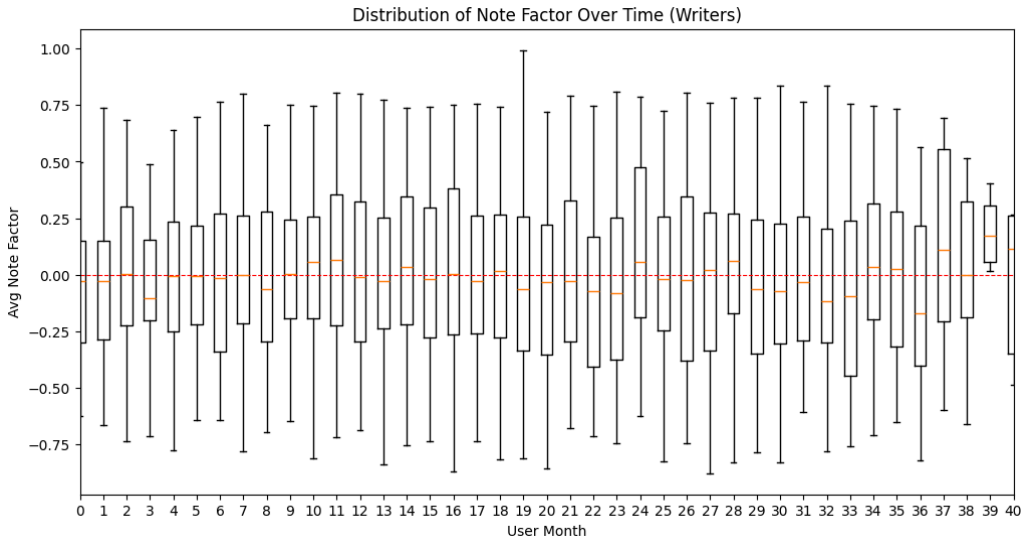

Looking at the distribution, it also remains stable over time, with the media staying near zero throughout, similarly to the average note factor over time graph. However, the spread of the distribution is wide at every month, which suggests that while the average does not lean, there is variation between individual writers. Some writers consistently product positive-factor notes while others consistently produce negative-factor notes.

I sampled 10 writers randomly for the purpose of seeing individual patterns clearer. Some writers consistently produce positive-factor notes, others consistently produce negative-factor notes, and some fluctuate heavily across months. This suggests that the neutral population average is not because writers are unbiased, but because writers with opposite political leanings cancel each other out.

For writers, also consider the average note intercept of notes they wrote, along with the number of CRH notes they wrote

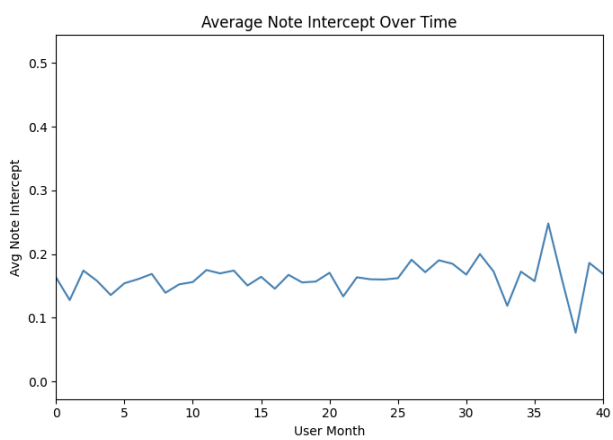

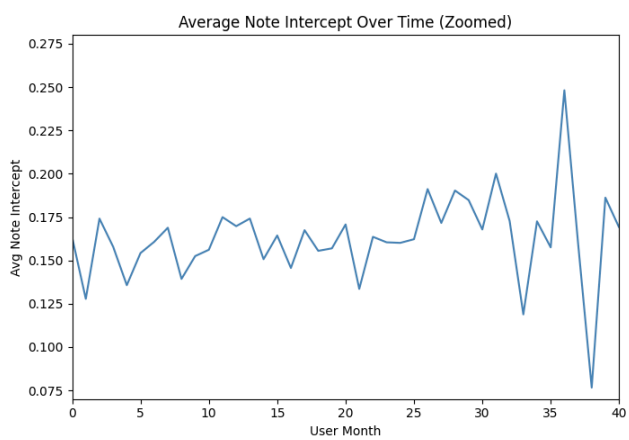

The average note intercept stays relatively stable around 0.15 throughout, with a slight upward trend after month 25. This suggests that notes maintain a consistent level of agreement across users with different political perspectives, rather than becoming more or less divisive. The modest increase in later months may indicate that more experienced writers produce notes that are slightly more broadly agreeable, although this pattern is weak and should be interpreted with caution due to noise in the data.

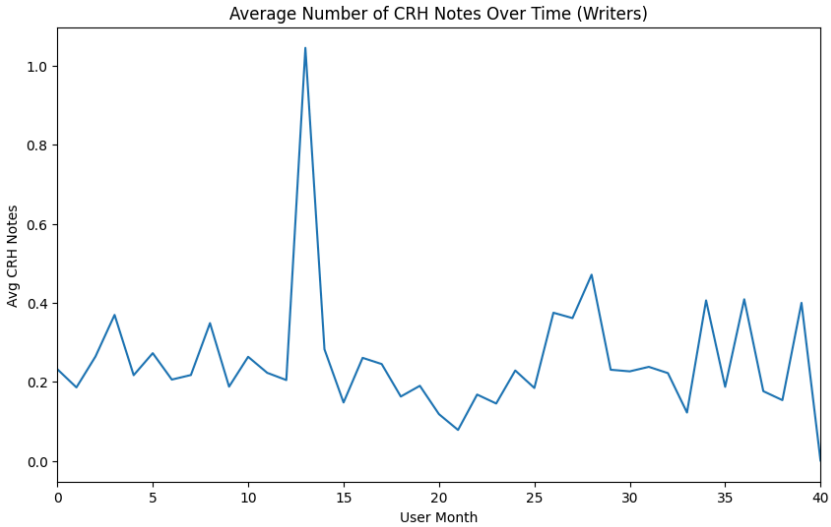

The average number of CRH notes per writer stays relatively low and stable around 0.2 to 0.3 throughout most months, with no clear upward trend. There is a notable spike at user month 13 which is likely driven by a small number of highly active writers at that point in their lifetime rather than a meaningful behavioral shift. Overall, writers do not appear to produce more CRH notes as they gain experience.

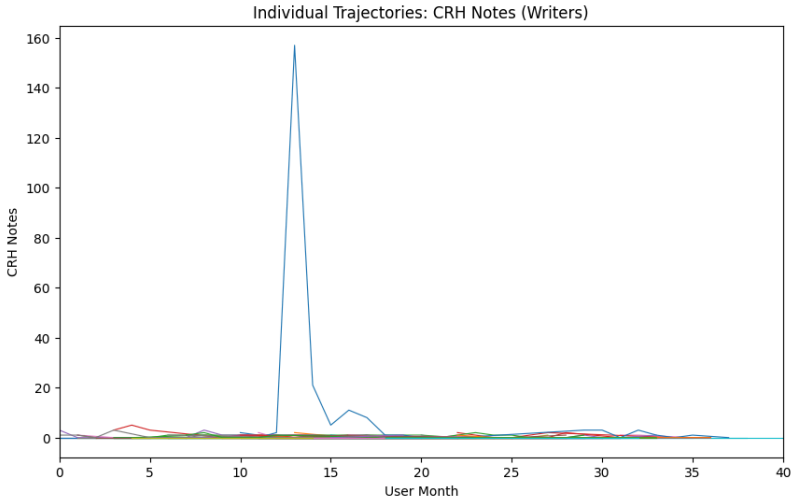

Looking at the individual trajectories with 100 sampled users, the spike at month 13 was caused by a single highly active writer who wrote an unusually large number of notes that month.

So overall, writers do not appear to improve over time. Their note intercept stays flat, and they do not produce more CRH notes as they gain experience. The notes they write maintain a consistent level of quality throughout their lifetime, suggesting that writing behavior is established early and does not meaningfully change.

Requesters

For requesters, consider the number of requests they made that resulted in notes and CRH notes.

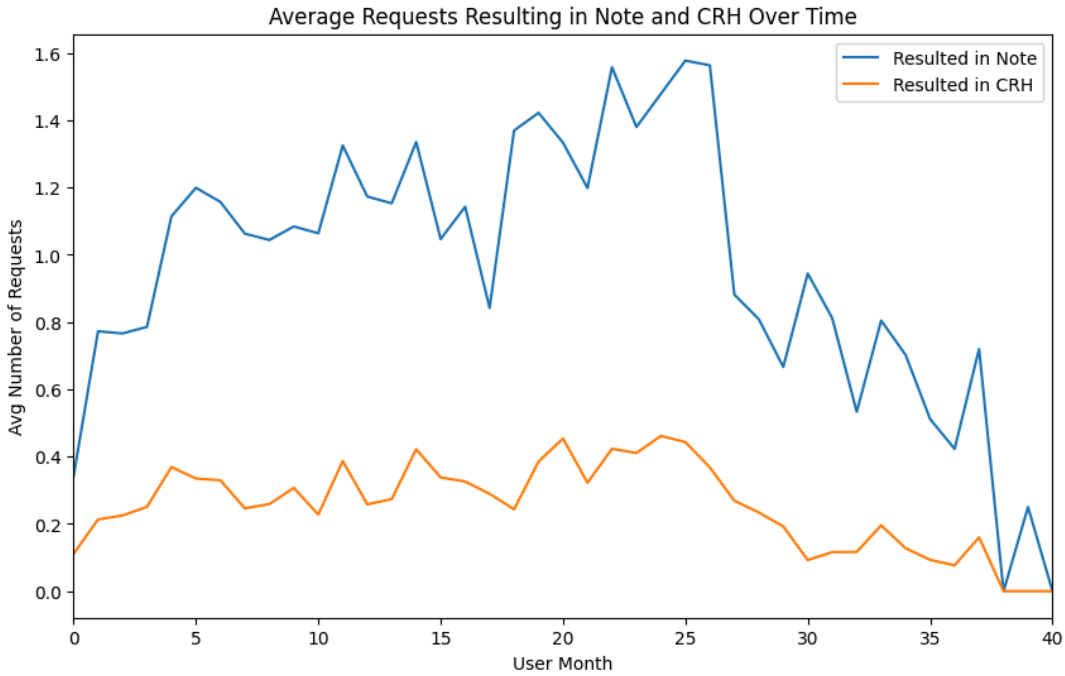

The number of requests resulting in a note increases gradually at month 0 to about 1.5 by month 25, suggesting that more experienced requesters get better at submitting requests that attract writers. However, both lines drop sharply after month 25. The number of requests resulting in a CRH note stays consistently lower and flatter throughout, meaning that while experienced requesters can attract writers, getting those notes to achieve CRH status does not improve as much over time.